

Articulations

- Definition:
 - A node v of a connected graph G is an **articulation** node (also called a cut vertex) if the removal of v and all its incident edges causes G to become disconnected.
- Motivation for articulations:
 - Articulations are important in communication networks.
 - In traffic flows they identify places that will stop traffic between two areas of a city if they become blocked.

Finding Articulations (1)

- Problem:
 - Given any graph $G = (V, E)$, find all the articulation points.
- Possible strategy:
 - For all vertices v in V :
 - Remove v and its incident edges.
 - Test connectivity using a DFS.
 - Execution time: $\Theta(n(n + m))$.
 - Can we do better?

Finding Articulations (2)

- A DFS tree can be used to discover articulation points in $\Theta(n + m)$ time.
 - We start with a program that computes a DFS tree labeling the vertices with their discovery times.
 - We also compute a function called $low[v]$ that can be used to characterize each vertex as an articulation or non-articulation point.
 - The root of the DFS tree (the root has a $d[\]$ value of 1) will be treated as a special case:

Finding Articulations (3)


- The root of the DFS tree is an articulation point if and only if it has two children.
 - Suppose the root has two or more children.
 - Recall that the back edges never link the vertices in two different subtrees.
 - So, the subtrees are only linked through the root vertex and if it is removed we will get two or more connected components (i.e. the root is an articulation point).
 - Suppose the root is an articulation point.
 - This means that its removal would produce two or more connected components each previously connected to this root vertex.
 - So, the root has two or more children.

Finding Articulations (4)

- Computation of $low[v]$.
 - We need another function defined on vertices:
 - This quantity will be used in our articulation finding algorithm:
$$low[v] = \min\{d[v], d[z] \text{ such that } (u, z) \text{ is a back edge for some descendent } u \text{ of } v\}$$
 - So, $low[v]$ is the discovery time of the vertex closest to the root and reachable from v by following zero or more edges downward, and then at most one back edge.


Finding Articulations (5)

- For non-root vertices we have a different test.
 - Suppose v is a non-root vertex of the DFS tree T . Then v is an articulation point of G if and only if there is a child w of v in T with $d[v] \leq low[w]$.
 - Sufficiency: Assume such a child w exists.
 - There is no descendent vertex of v that has a back edge going "above" vertex v .
 - Also, there is no cross link from a descendent of v to any other subtree.
 - So, when v is removed the subtree with w as its root will be disconnected from the rest of the graph.



Finding Articulations (6)

- Necessity: Assume no such child w exists.
 - In this case all children of v have a descendent with a back edge going to an ancestor of v .
 - When v is removed each of the children of v will still be connected to some vertex on the path going from the root to the vertex v .
 - The graph stays connected and so v would not be an articulation point in this case.



Finding Articulation Points Pseudocode

```

function dfs-visit(v)
  status[v] := gray; time := time+1; d[v] := time;
  low[v] := d[v];
  for each w in out(v)
    if status[w] == white
      // In this case (v,w) is a TREE edge
      dfs-visit(w);    // NOTE: low[w] is now computed!
      if d[v] <= low[w] then
        record that vertex v is an articulation
      low[v] := min(low[v], low[w]) // Note how low[ ] can propagate up to a parent.
    else if w is not the parent of v then
      // In this case (v,w) is a BACK edge
      low[v] := min(low[v], d[w])
  status[v] := black;
  
```

Minimum Spanning Trees

- Problem:
 - Given a connected undirected weighted graph $G = (V, E)$, find a minimum spanning tree T for G .
- Assumptions
 - Weights are nonnegative.
 - The cost of a spanning tree is the sum of all the weights of all the edges in T .
 - The Minimum Spanning Tree (MST) is the spanning tree with the smallest possible cost.
- Typical application: Connect nodes in a computer network using as little wire as possible (MST links).

Kruskal's Algorithm

```
// Sort edges in order of increasing weight
// so that  $w[f[1]] \leq w[f[2]] \leq \dots \leq w[f[m]]$ 

T := empty set;
for i:=1 to m do
    let u,v be the endpoints of edge f[i]
    if there is no path between u and v in T then
        add f[i] to T
return T
```

Correctness of Kruskal's Algorithm

- Kruskal's algorithm produces a MST:
 - Kruskal's greedy algorithm produces a tree T_G .
Let edges be e_1, e_2, \dots, e_{n-1} sorted by weight.
 - Then for any $0 \leq k \leq n - 1$ there exists a minimum spanning tree that contains edges e_1, e_2, \dots, e_k .
- Proof by induction:
 - Base case:
 - For $k = 0$ the lemma holds trivially.
 - Induction step:

Correctness of Kruskal's Algorithm

- Suppose there is a MST T^* with edges: e_1, e_2, \dots, e_{k-1} .
- **Case 1:** $e_k \in T^*$:
 - Then T^* contains all the edges e_1, e_2, \dots, e_k and the statement is true.
- **Case 2:** $e_k \notin T^*$:
 - If we remove e_k from T_G , then T_G becomes disconnected and will have two components (call them A and B).
 - Add e_k to T^* . This creates a cycle in T^* involving vertices in both A and B .
 - So, the cycle must contain an edge e' different from e_k that has one endpoint in A and one in B .
 - Remove edge e' , to obtain a new graph T' $\implies T'$ is a spanning tree.

Correctness of Kruskal's Algorithm

- Note that $w(e') \geq w(e_k)$, otherwise e' would have been chosen by Kruskal's algorithm instead of e_k .
- The cost of T' can be written as:
$$w(T') = w(T^*) + w(e_k) - w(e') \quad \text{implying}$$
$$w(T') \leq w(T^*).$$
- Since T^* is a MST, $w(T') = w(T^*)$ and T' is also a MST.
- Moreover, T' contains each of the edges e_1, e_2, \dots, e_k which is what we wanted to prove.
- Thus, we have proved by induction that for every k there exists a MST that contains each of the edges e_1, e_2, \dots, e_k .

Analysis of Kruskal's Algorithm

- Running time:
 - Sorting the edges takes $\Theta(m \log m) = \Theta(m \log n)$ time.
 - Running time for the rest of algorithm depends on implementation of the path detection statement:
"if there is no path between u and v in T "
 - Use DFS on the edges of T selected so far:
 - There are less than n of them, so it will take $O(n)$ per check.
 - This implies a final running time that is $O(mn)$.
 - Use a Union/Find data structure (covered in CS466):
 - The check would take $O(\log n)$ (or better) for each check.
 - This implies a final running time that is $O(m \log n)$.

Prim's Algorithm

- Main idea:
 - Start from an arbitrary single vertex s and gradually “grow” a tree.
 - We maintain a set of connected vertices S .

```
S := {s};  
T := empty set;  
while S <> V do  
    e := (u,v) such that u is in S, v is not  
        in S and w(e) is smallest possible;  
    add v to S;  
    add e to T;  
return T;
```

Correctness of Prim's Algorithm

- Prim's algorithm produces a MST:
 - Let Prim's greedy algorithm produce a tree T_G containing edges: e_1, e_2, \dots, e_{n-1} (numbered in the order they were added by the algorithm).
 - Then for any $0 \leq k \leq n - 1$ there exists a minimum spanning tree that contains edges e_1, e_2, \dots, e_k .
- Proof by induction:
 - Base case:
 - For $k = 0$ the lemma holds trivially.
 - Induction step:

Correctness of Prim's Algorithm

- Suppose there is a MST T^* with edges: e_1, e_2, \dots, e_{k-1} .
- **Case 1:** $e_k \in T^*$:
 - Then T^* contains all the edges e_1, e_2, \dots, e_k and the statement is true.
- **Case 2:** $e_k \notin T^*$:
 - Let S be the set of finished vertices after $k - 1$ steps of the algorithm.
 - Add e_k to T^* . This will create a cycle in T^* .
 - The cycle must contain an edge e' different from e_k that has one endpoint in S and one not in S .
 - Remove edge e' and denote the new graph by T' .
 - T' is a spanning tree.

Correctness of Prim's Algorithm

- Note that $w(e') \geq w(e_k)$, otherwise e' would have been chosen by Prim's algorithm instead of e_k .
- The cost of T' can be written as:
$$w(T') = w(T^*) + w(e_k) - w(e') \quad \text{implying}$$
$$w(T') \leq w(T^*).$$
- Since T^* is a MST, $w(T') = w(T^*)$ and T' is also a MST.
- Moreover, T' contains each of the edges e_1, e_2, \dots, e_k which is what we wanted to prove.
- Thus, we have proved by induction that for every k there exists a MST that contains each of the edges e_1, e_2, \dots, e_k .

Analysis of Prim's Algorithm

- Running time:
 - We can improve the algorithm by keeping for each vertex not in S its least cost neighbour in S .
 - The cost for this neighbour will be stored in $cost[v]$ and the neighbour itself in $other[v]$. (See next page).
 - We do the same set of operations with the cost as in Dijkstra's algorithm:
(initialize a structure, decrease values m times, select the minimum $n - 1$ times).
 - Therefore we get $O(n^2)$ time when we implement cost with an array, and $O((n + m) \log n)$ when we implement it with a heap.

Pseudocode for Prim's Algorithm

```
S := {s};
T := empty set;
// Initialize data structure
for each u not in S
    cost[u] := w(s,u);
    other[u] := s;
// Main computation
while S <> V do
    v := vertex not in S with the smallest cost[v];
    e := (v, other[v]);
    add v to S;
    add e to T;
    // Update data structure
    for each x not in S and a neighbour of v
        if w(v,x) < cost[x] then
            cost[x] := w(v,x);
            other[x] := v;
return T;
```

Recall: $cost[x]$ is least cost between x and "nearest" vertex in S .

Note: $w(u,v)$ is defined for all possible u and v in V . When there is **no** edge between u and v we have $w = \text{infinity}$.

Dijkstra's Algorithm

- Objective of Dijkstra's algorithm:
 - Dijkstra's algorithm finds the least cost paths from a source vertex s to **all** the other vertices in the graph.

Dijkstra's Algorithm Setup

- We maintain 2 sets of vertices:
 - **The set C of “finished” vertices.**
 - We can think of C as the “cloud set”.
 - It will start with the source vertex and eventually expand to include all the other vertices.
 - For any vertex in the cloud we will be assured that we know its least cost path to the source.
 - **The set Q of vertices that are yet to be processed.**
 - They are not in the cloud.


Dijkstra's Algorithm: Pseudocode

```
function least_cost(s)
  // Initialize costs
  C :=  $\emptyset$ ; Q := V;
  for all w in V do cost[w] := infinity;
  cost[s] := 0;
  // Iterative steps:
  while Q is non-empty do           //Note: Greedy!
    v := vertex such that cost[v] is a minimum;
    move vertex v from Q to the cloud set C;
    // Update the costs
    for all t in Q that are in out(v) do
      cost_via_v := cost[v] + c[v, t];
      if(cost_via_v < cost[t]) then
        cost[t] := cost_via_v;
```

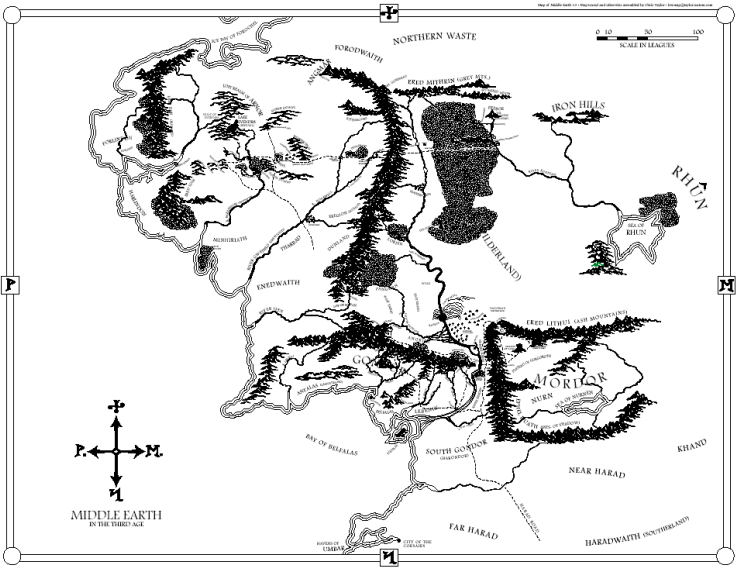
- Path reconstruction:
 - We keep the last but one vertex in the shortest path.

An Example of Dijkstra's Algorithm

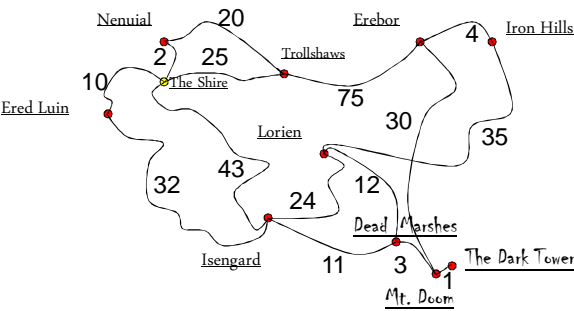
- Suppose small people must minimize travel costs because they walk around with big hairy feet and no shoes....



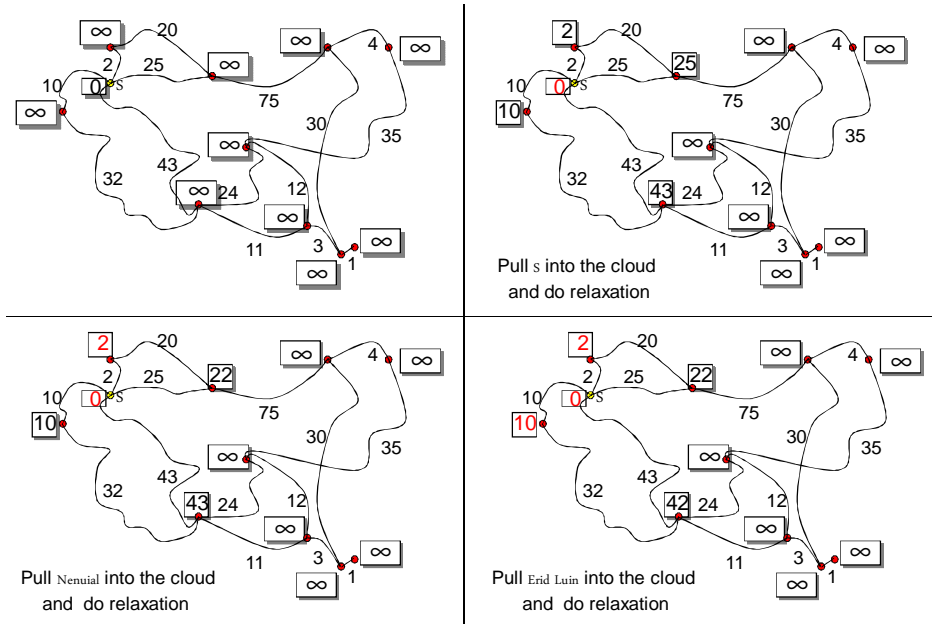
Middle Earth (In the Third Age)



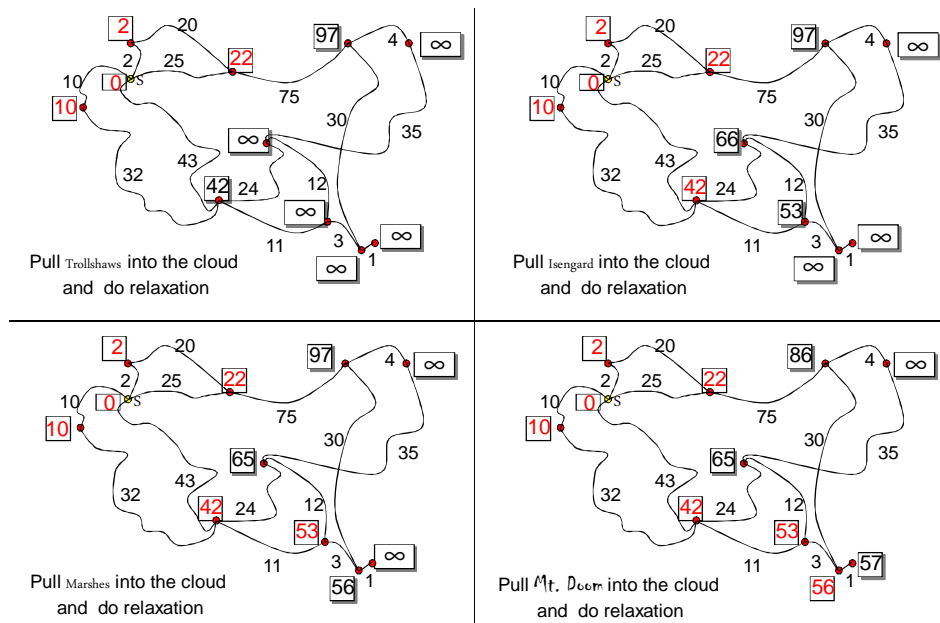
Middle Earth: Travel Cost



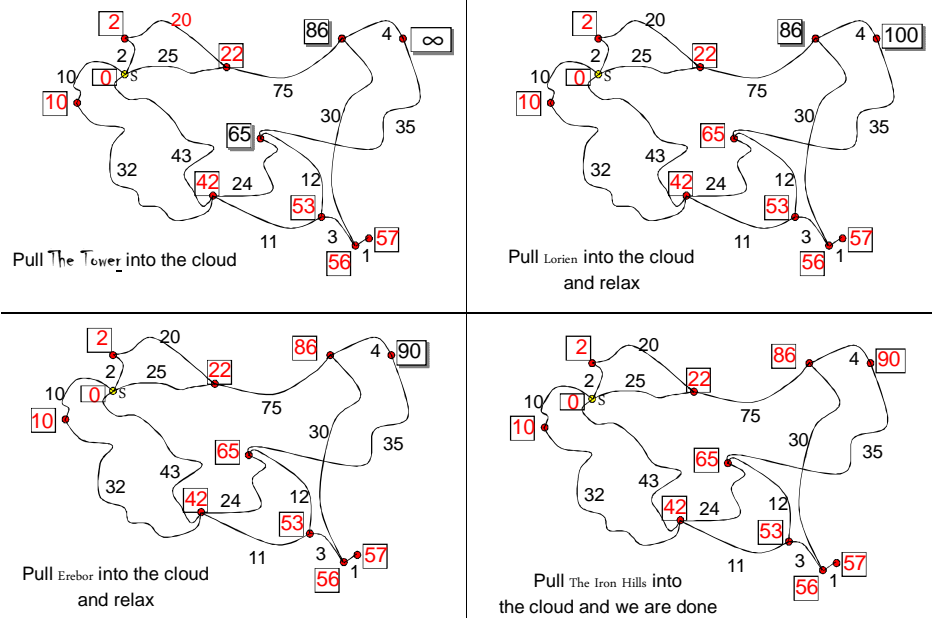
Dijkstra Clouds Middle Earth



Dijkstra Clouds Middle Earth



Dijkstra Clouds Middle Earth



Running Time of Dijkstra's Algorithm

- Running time will depend on the implementation of the data structure for $cost[s]$.

The Floyd-Warshall Algorithm

- Problem:
 - Given a graph $G = (V, E)$, directed or undirected, weighted with edge costs, find the least cost path from u to v for all pairs of vertices (u, v) .
 - We assume all weights are **non-negative numbers**.
 - The cost of a path will be the sum of the costs of all edges in the path.

Floyd-Warshall: a Useful Lemma

- Lemma:
 - Let P be the least cost path from u to v .
 - Consider any two vertices x and y on this path.
 - The part of the path between vertices x and y will be the least cost path between x and y .

Proof:

- If there was a subpath from x to y that was not the least cost path from x to y , then we could replace this subpath with the least cost path from x to y , obtaining a lesser cost for the overall path.
- This contradicts our statement that the path from u to v was the shortest path, so the lemma is true.

Floyd-Warshall: Extending the Cost Function

- The previous lemma suggest the possibility of using a dynamic programming strategy for our problem.
- A useful way to look at the problem:
 - It is convenient to think of the problem as having a cost $c(u, v)$ assigned to each of the pairs for **all** possible pairs u and v in the graph.
 - $c(u, v) =$ the given edge cost if edge (u, v) exists.
 - $c(u, v) =$ infinity if there is no edge (u, v) in the graph.

Floyd-Warshall: Extending the Cost Function

- With the extended definition of cost, we can go from u to v using any subset of distinct vertices (apart from u and v) as intermediate nodes in the path.
 - Of course, if the selected path uses a non-existent edge in G , the cost of the path is infinity.
 - The algorithm will discard paths with infinite cost and so we will get solutions made up from the given edges.
 - So, the algorithm will examine all possible paths without the need to check beforehand if edges actually exist in G .

Floyd-Warshall: Subproblem Definition

- Subproblem setup:
 - We assume the vertices are labeled (i.e. indexed) using integers ranging from 1 to n .
 - An adjacency matrix representation of the graph is convenient.
- Subproblem definition:
 - We let $cost[i, j, k]$ hold the cost of the least cost path between vertex i and vertex j with intermediate nodes chosen from vertices $1, 2, \dots, k$.

Floyd-Warshall: Subproblem Definition

- Recall our subproblem definition:
 - We let $cost[i, j, k]$ hold the cost of the least cost path between vertex i and vertex j with intermediate nodes chosen from vertices $1, 2, \dots, k$.
 - As the index k increases we have more options for discovering the shortest path between endpoints i and j .
 - Even if there is an edge from i to j , its cost might exceed that of another path running from i to j .
 - So, the least cost for the path from i to j will be $cost[i, j, n]$, that is, we have the option of selecting from **all** the other nodes different from i and j .
 - Base case: $cost[i, j, 0] = c(i, j)$. ← (Given edge costs)
 - $cost[i, j, 0]$ is for the path with **no** intermediate nodes.

Floyd-Warshall: The Recurrence

- How do we evaluate $cost[i, j, k]$?
 - Our strategy will be to evaluate all $cost[]$ values starting with $k = 1$, then $k = 2$, etc.
 - Recall that the least cost path for $cost[i, j, k]$ can involve any intermediate nodes selected from $\{1, 2, \dots, k\}$.
 - In particular, the least cost path may involve node k or it may not...
 - **Case 1**: The least cost path does **not** go through node k , then $cost[i, j, k] = cost[i, j, k-1]$.
 - **Case 2**: The least cost path **does** go through node k , then $cost[i, j, k] = cost[i, k, k-1] + cost[k, j, k-1]$.

Floyd-Warshall: The Recurrence

- Of course, we want to use the case that gives us the smaller cost:

$$cost[i, j, k] = \min\{cost[i, j, k-1], cost[i, k, k-1] + cost[k, j, k-1]\}$$

Some improvements:

- The value of $cost[i, j, k]$ is always dependent on the immediately previous cost values corresponding to the third parameter equal to $k-1$ (i.e. not dependent on $k-2, k-3$, etc.)
- So, we can do away with the third parameter and keep the costs in a two dimensional array that is updated n times.
- Thus, $cost[i, j, k]$ will remain as $cost[i, j, k-1]$ unless we update it with a smaller $cost[i, k, k-1] + cost[k, j, k-1]$ value.

Floyd-Warshall: Pseudocode

```
for i := 1 to n do
  for j := 1 to n do
    cost[i, j] := c[i, j]; // Let c[u, u] := 0
for k := 1 to n do
  for i := 1 to n do
    for j := 1 to n do
      sum = cost[i, k] + cost[k, j];
      if(sum < cost[i, j]) then cost[i, j] := sum;
```

- This code derives the least cost value but there is no recovery of the actual path.
- This is done by remembering the second vertex of the path found so far:

Floyd-Warshall: Pseudocode

```
for i := 1 to n do
  for j := 1 to n do
    cost[i, j] := c[i, j];
    next[i, j] := j; // Note!
for k := 1 to n do
  for i := 1 to n do
    for j := 1 to n do
      sum := cost[i, k] + cost[k, j];
      if(sum < cost[i, j]) then
        cost[i, j] := sum;
        next[i, j] := next[i, k]; // Note!
// To write out the path from u to v:
w := u;
write w;
while w != v do
  w := next[w, v];
  write w;
```

Note: Running time $\in \Theta(n^3)$.

Formulating Problems as Graph Problems

- As a review we now look at four problems.
 - You should read the problems and as homework try to solve them without looking at the answers in the slides that follow.

Formulating Problems as Graph Problems: Problem #1

- Reliable network routing:
 - Suppose we have a computer network with many links.
 - Every link has an assigned reliability.
 - The reliability is a probability between 0 and 1 that the link will operate correctly.
 - Given nodes u and v , we want to choose a route between nodes u and v with the highest reliability.
 - The reliability of a route is a product of the reliabilities of all its links.

Problem #2


- The Greyhound bus problem:
 - Suppose we are given a bus schedule with information for several buses. A bus is characterized by four attributes:
 - the “from-city”, the “to-city”, departure time, arrival time.
 - Find buses going from city F to city T taking the fastest trip?
 - Take into account travel and wait times between bus arrivals and depatures..
 - First, we eliminate an idea that leads to an inadequate solution:
 - Use a graph that has nodes representing cities.
 - Label each edge with the travel time between cities.
 - Now go for the least cost path.
 - BUT: there is no accounting for wait times!
 - Also, travel times between two cities may vary during the day.
 - But there is another way to use a graph strategy...

Sample Bus Schedule

UW to Hamilton	15:40 17:25		
UW to Toronto	09:00 11:00	17:00 19:00	
Hamilton to Niagara Falls	17:30 18:45		
Toronto to Niagara Falls	12:30 14:05	20:30 22:10	
Niagara Falls to Buffalo	14:10 15:25	18:40 19:40	22:55 23:59

Problems #3

- The RootBear Problem:
 - Suppose we have a narrow canyon with perpendicular walls on either side of a forest.
 - We assume a north wall and a south wall.
 - Viewed from above we see the A&W RootBear attempting to get through the forest.
 - We assume trees are represented by points.
 - We assume the bear is a circle of given diameter d .
 - We are given a list of coordinates for the trees.
 - Find an algorithm that determines whether the bear can get through the forest.



Solution to Problem #1

- Reliable network routing:
 - Suppose we have a computer network with many links.
 - Every link has an assigned reliability.
 - The reliability is a probability between 0 and 1 that the link will operate correctly.
 - Given nodes u and v , we want to choose a route between nodes u and v with the highest reliability.
 - The reliability of a route is a product of the reliabilities of all its links.

- The route will correspond to a path in the graph.
- Can we make this look like a shortest path problem?
- Yes:
 - Since reliability is computed as a product, we will want to change the weights so that an edge is assigned the logarithm of the probability.
 - Then we sum logs to work with products of probabilities.
 - To get the best reliability path we want the highest probability of operation which we can derive by finding the least weight path if the assigned weights are *negative* logarithms of the probability values.
 - Then we are able to use Dijkstra's algorithm.

Solution to Problem #2

- The Greyhound bus problem:
 - Suppose we are given a bus schedule with information for several buses. A bus is characterized by four attributes:
 - the “from-city”, the “to-city”, departure time, arrival time.
 - Find buses going from city F to city T with the fastest trip?
 - Take into account travel and wait times between arrival and departure times..
 - First, let's eliminate an idea leading to an inadequate solution:
 - Use a graph that has nodes representing cities.
 - Label each edge with the travel time between cities.
 - Now go for the least cost path.
 - BUT: there is no accounting for wait times!
 - Also, travel times between two cities may vary during the day.
 - But there is another way to use a graph strategy...

Sample Bus Schedule

UW to Hamilton	15:40 17:25		
UW to Toronto	09:00 11:00	17:00 19:00	
Hamilton to Niagara Falls	17:30 18:45		
Toronto to Niagara Falls	12:30 14:05	20:30 22:10	
Niagara Falls to Buffalo	14:10 15:25	18:40 19:40	22:55 23:59


- Another approach:
 - Use a graph in which each vertex is a bus.
 - There will be an edge between busses x and y if and only if:

$$x.to_city = y.from_city \quad \text{and} \\ y.departure_time \geq x.arrival_time.$$

- Our time cost for an edge will be:

$$\begin{aligned} & \text{waiting time} + \text{travel time on bus } y = \\ & (y.departure_time - x.arrival_time) + \\ & \quad (y.arrival_time - y.departure_time) \\ & = y.arrival_time - x.arrival_time. \end{aligned}$$


- We need two special vertices for the origin and destination cities.
- There is an edge from origin to bus x if and only if $x.from_city = origin$.
- Time cost of this edge is $x.arrival_time - x.departure_time$.
- There is an edge from bus y to the destination if and only if $y.to_city = destination$.
- The time cost of this edge is 0.
- We now have a shortest path problem:



- Note: the shortest trip is via Toronto with time 6:25 hours.

Solution to Problem #3

- The RootBear Problem:
 - Suppose we have a canyon with perpendicular walls on either side of a forest.
 - We assume a north wall and a south wall.
 - Viewed from above we see the A&W RootBear attempting to get through the forest.
 - We assume trees are represented by points.
 - We assume the bear is a circle of given diameter d .
 - We are given a list of coordinates for the trees.
 - Find an algorithm that determines whether the bear can get through the forest.



- The graph formulation for this problem:
 - Create a vertex for each tree, and a vertex for each canyon wall.
 - Two trees are connected by an edge if and only if the RootBear cannot pass between them.
 - That is if their separation is less than d .
 - Do the same for a tree and its perpendicular distance to a canyon wall.
 - Now determine if canyon walls are in the same connected component of the graph.
 - If they are then the bear cannot pass through the canyon.
 - Otherwise the boundary of the connected component containing the northern canyon wall defines a viable path for the bear.



Conclusion

- Graphs are a very important formalism in computer science.
- Efficient algorithms are available for many important problems:
 - exploration, shortest paths, minimum spanning trees, cut links, colouring, etc.
- If we formulate a problem as a graph problem, chances are that an efficient non-trivial algorithm for solving the problem is already known.
- Some problems have a natural graph formulation.
 - For others we need to choose a less intuitive graph formulation.
- Some problems that do not seem to be graph problems at all can be formulated as such.